

Simplemux v2 Readme file

Content

About Simplemux	2
Modes.....	3
Flavors.....	3
Tunneling modes	3
Header compression	5
Tunneling	5
Multiplexing	6
What will you find in this package	12
Required tools	12
ROHC	12
Compiling simplemux	12
Usage of simplemux	12
If you have to use the same local interface more than once	13
Format of the Simplemux traces	13
Trace examples	15
Scripts for calculating compression statistics	15
Scripts for drawing the instantaneous throughput and pps	16
Scenario 1: two Raspberries.....	18
Scenario 2: three machines	20
Create a tun interface in machine 4.....	20
Create a tun interface in machine 5.....	20
Establish the simplemux tunnel between machine 4 and machine 5	20
Steer the desired traffic to the tunnel using <i>ip rule</i> and <i>iptables</i>	21
Examples implementing different policies.....	21
Credits.....	23

About Simplemux

There are some situations in which multiplexing a number of small packets into a bigger one is desirable. For example, a number of small packets can be sent together between a pair of machines if they share a common network path. Thus, the traffic profile can be shifted from small to larger packets, reducing the network overhead and the number of packets per second to be managed by intermediate routers.

Simplemux is a generic multiplexing protocol, described in [draft-saldana-tsvwg-simplemux](#). It is able to encapsulate a number of packets belonging to different protocols into a single packet. It includes the "Protocol" field on each multiplexing header, thus allowing the inclusion of a number of packets belonging to different protocols in a packet of another protocol (Fig. 1). The size of the multiplexing headers is kept very low (it may be a single byte when multiplexing small packets) in order to reduce the overhead.




Fig. 1. Example of a Simplemux packet including packets of different protocols

Simplemux is designed to optimize together several flows sharing a common network path or segment (Fig. 2). Optimization in the end host is (in principle) not useful. The multiplexing is performed between a pair of machines called ingress-optimizer and egress-optimizer.




Fig. 2. Scheme of the optimization

The present code is an implementation of **Simplemux**, using the next options for the *tunneling protocol*:

- IP
- UDP/IP
- TCP/IP

and the next options for the *multiplexed protocol*:

- IP packets
- ROHC ([RFC 5795](#)) compressed packets
- Ethernet frames

As illustrated in Fig. 3, it can perform traffic optimization, combining multiplexing with header compression. This may result on significant bandwidth savings and pps reductions for small-packet flows (e.g. VoIP, online games).

Native traffic: Five IPv4/UDP/RTP VoIP packets with two samples of 10 bytes




Optimized traffic (network mode): One IPv4 simplemux Packet including **five** RTP packets



Optimized traffic (transport mode): One IPv4 simplemux Packet including **five** RTP packets



Native traffic headers:



Optimized traffic headers:




Fig. 3. Optimization of five VoIP RTP packets

Modes

Simplemux can run in four *modes*:

- **Network mode:** the multiplexed packet is sent in an IP datagram using Protocol Number 253 or 254 (according to IANA, these numbers can be used for experimentation and testing¹).
- **UDP mode:** the multiplexed packet is sent in an UDP/IP datagram. In this case, the protocol number in the outer IP header is that of UDP (17) and both ends must agree on a UDP port (the implementation uses 55555 or 55557 by default).
- **TCP server mode:** the multiplexed packet is sent in a TCP/IP datagram. In this case, the protocol number in the outer IP header is that of TCP (4) and both ends must agree on a TCP port (the implementation uses 55555 or 55557 by default).
- **TCP client mode:** same as TCP server mode, but it is the TCP client.

Flavors

Simplemux has two *flavors*:

- **Normal:** it tries to compress the separators as much as possible. For that aim, some single-bit fields are used.
- **Fast:** it sacrifices some compression on behalf of speed. In this case, all the separators are 3-byte long, and all have the same structure.

Simplemux *fast* must be used in TCP mode.

Tunneling modes

Simplemux uses Linux TUN/TAP virtual interface. Two *tunneling modes* are allowed:

- **TUN tunneling mode:** IP packets are multiplexed between the two endpoints.
- **TAP tunneling mode:** Ethernet frames are multiplexed between the two endpoints.

Fig. 4 shows the protocol stack used for sending an Ethernet frame (TAP tunneling mode) carrying a UDP/IP packet, using IP, UDP and TCP modes.

¹ Protocol numbers, <http://www.iana.org/assignments/protocol-numbers/protocol-numbers.xhtml>

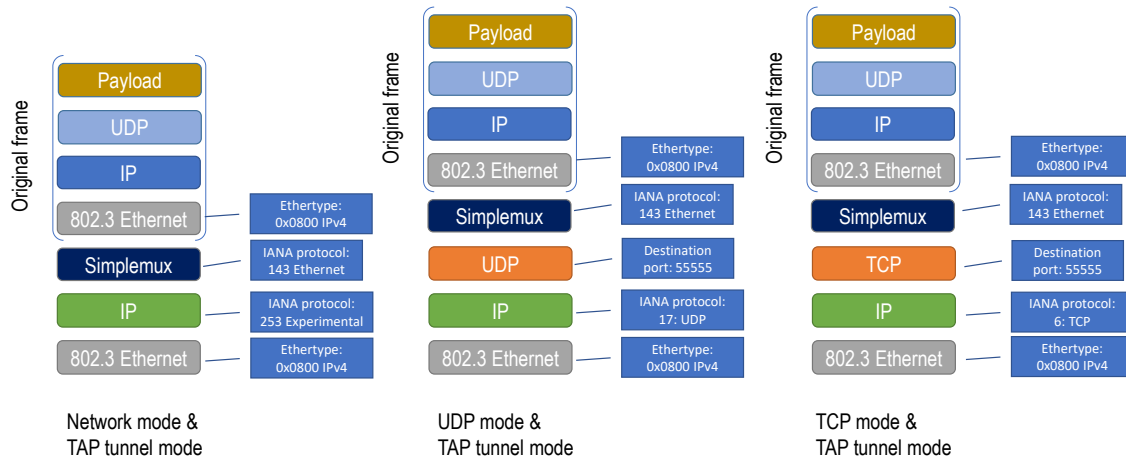


Fig. 4. Sending a UDP Ethernet frame in TAP tunneling mode: IP, UDP and TCP modes.

Fig. 5 presents the protocol stack used for sending a UDP/IP uncompressed packet (TUN mode), using IP, UDP and TCP modes.

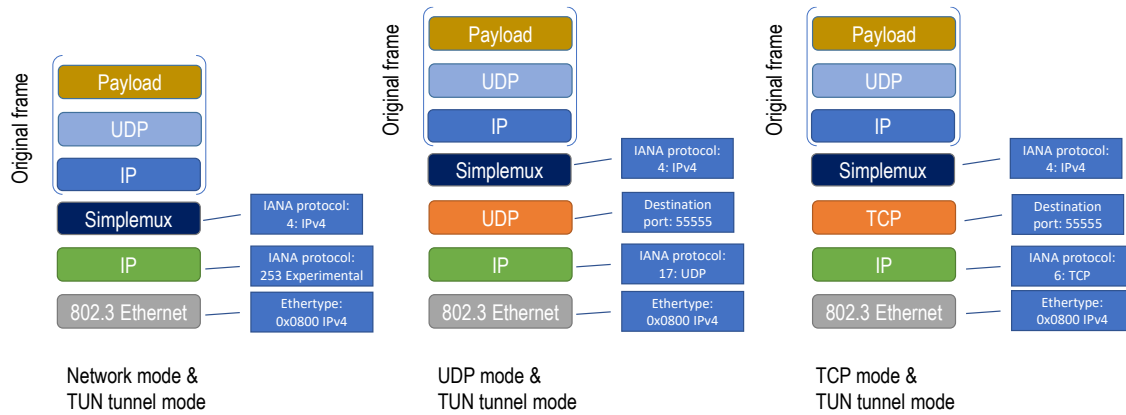


Fig. 5. Tunneling a UDP packet in TUN tunneling mode: IP, UDP and TCP modes.

Fig. 6 presents the protocol stack used for sending a ROHC compressed packet (TUN mode), using IP, UDP and TCP modes.

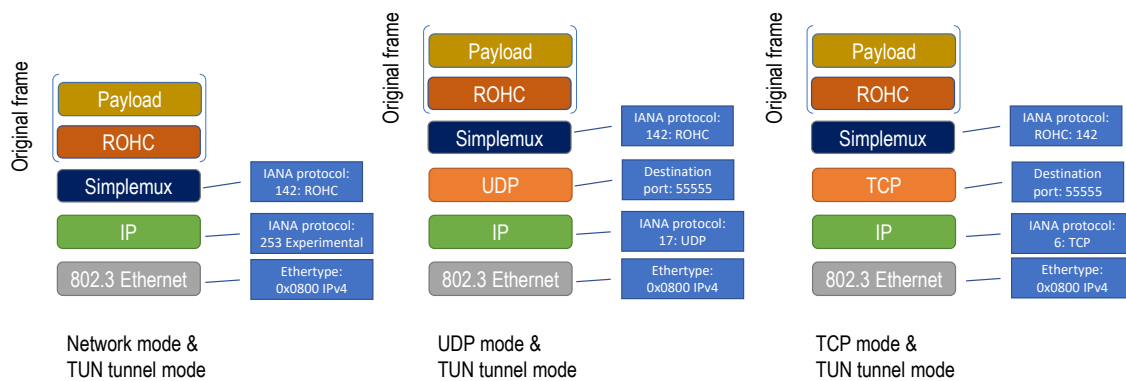


Fig. 6. Tunneling a ROHC compressed packet in TUN tunneling mode: IP, UDP and TCP modes.

Simplemux can be used as an option at the multiplexing layer of TCM², a combination of protocols for Tunneling, Compressing and Multiplexing, allowing the optimization of small-packet flows. TCM may use of a number of different standard algorithms for header compression, multiplexing and tunneling, combined in a similar way to [RFC 4170](#), as proposed in [draft-saldana-tsvwg-tcmf](#).

The implementation of Simplemux is written in C for Linux. It compresses headers using an implementation of ROHC by Didier Barvaux (<https://rohc-lib.org/>).

Header compression

This **Simplemux** implementation includes these ROHC modes:

- ROHC unidirectional.
- ROHC bidirectional optimistic
- ROHC bidirectional reliable is not yet implemented.

ROHC cannot be enabled in one of the peers and disabled in the other peer.

ROHC is able to compress this traffic flows:

- IP/UDP/RTP: If the UDP packets have the destination ports 1234, 36780, 33238, 5020, 5002, the compressor assumes that they are RTP.
- IP/UDP
- IP/TCP
- IP/ESP
- IP/UDP-Lite

Tunneling

In this implementation, tunneling is performed this way:

- In **network mode**, the external IP header is the tunneling header. It uses Protocol number 253 or 254 (*fast flavor*).
- In **UDP mode**, a UDP/IP header is added before the first **Mux separator**. By default, the destination UDP port is 55555 or 55557 (*fast flavor*).
- In **TCP mode**, a TCP/IP header is added before the first Mux separator. By default, the destination TCP port is 55555 or 55557 (*fast flavor*).

ROHC feedback information (when using ROHC Bidirectional modes) is sent in UDP packets using port 55556 by default.

² Jose Saldana *et al*, "[Emerging Real-time Services: Optimizing Traffic by Smart Cooperation in the Network](#)," [IEEE Communications Magazine](#), Vol. 51, n. 11, pp 127-136, Nov. 2013.
[doi 10.1109/MCOM.2013.6658664](#)

Multiplexing

Normal flavor

The **Simplemux separator** (see Fig. 7) has two different formats: one for the *First header* (the separator before the first packet included in the multiplexed bundle), and another one for *Non-first headers* (the rest of the separators).




Fig. 7. First and Non-first Simplemux headers (also known as separators)


Format of the First Simplemux header/separator

In order to allow the multiplexing of packets of any length, the number of bytes expressing the length is variable, and the field Length Extension (**LXT**, one bit) is set to **0** if the current byte is the last one including length information.


These are the fields of the header:

- **Single Protocol Bit (SPB, one bit)** only appears in the first Simplemux header. It is set to **1** if all the multiplexed packets belong to the same protocol (in this case, the *Protocol* field will only appear in the first Simplemux header). It is set to **0** when each packet MAY belong to a different protocol.
- **Length Extension (LXT, one bit)** is **0** if the current byte is the last byte where the length of the first packet is included, and **1** in other case.
- **Length (LEN, 6, 13, 20, etc. bits)**. This is the length of the multiplexed packet (in bytes), not including the length field. If the length of the multiplexed packet is less than 64 bytes (less than or equal to 63 bytes), the first **LXT** is set to **0** and the 6 bits of the length field are the length of the multiplexed packet. If the length of the multiplexed packet is equal or greater than 64 bytes, additional bytes are added. The first bit of each of the added bytes is the **LXT**. If **LXT** is set to **1**, it means that there is an additional byte for expressing the length. This allows to multiplex packets of any length (see Fig. 8).
- **Protocol (8 bits)** is the *Protocol* field of the multiplexed packet, according to IANA "Assigned Internet Protocol Numbers".

packet length < 64 bytes



packet length ≥ 64 bytes
< 8192 bytes



packet length ≥ 8192 bytes
< 1048576 bytes




Fig. 8. Fields of the First Simplemux header/separator

For example, in the case of a packet of 65 bytes, the Simplemux separator will be 3 bytes long. In this case, the length of the packet will be the number expressed by the concatenation of the bits of Length 1 - Length 2 (total 13 bits). Length 1 includes the 6 most significant bits and Length 2 the 7 less significant bits.


More bytes can be added to the length if required, using the same scheme: 1 **LXT** byte plus 7 bits for expressing the length.

Format of the Non-first Simplemux header/separator


The Non-first Simplemux headers also employ a format allowing the multiplexing of packets of any length, so the number of bytes expressing the length is variable, and the field Length Extension (**LXT**, one bit) is set to 0 if the current byte is the last one including length information.

If the **SPB** (Single Protocol Bit) of the First Simplemux header is set to 1, it means that all the multiplexed packets belong to the same protocol. In this case, the format is the one presented in Fig. 9:

packet length < 128 bytes



packet length ≥ 128 bytes
< 16384 bytes (2^{14})



packet length ≥ 16384 bytes (2^{14})
< 2097152 bytes (2^{21})






Fig. 9. Fields of the Non-first Simplemux header/separator, when the **SPB** bit of the First header is 1

If the **SPB** of the First Simplemux header is set to 0, then each packet may belong to a different protocol, so the *Protocol* field is also included, as shown in Fig. 10:

packet length < 128 bytes



packet length ≥ 128 bytes
< 16384 bytes (2^{14})



packet length ≥ 16384 bytes (2^{14})
< 2097152 bytes (2^{21})




Fig. 10. Fields of the Non-first Simplemux header/separator, when the **SPB** bit of the First header is 0

These are the fields:

- **Length Extension (LXT, one bit)** is 0 if the current byte is the last byte where the length of the first packet is included, and 1 in other case.

- **Length (LEN, 7, 14, 21, etc. bits).** This is the length of the multiplexed packet (in bytes), not including the *Length* field. If the length of the multiplexed packet is less than 128 bytes (less than or equal to 127 bytes), the first **LXT** is set to 0 and the 7 bits of the length field are the length of the multiplexed packet. If the length of the multiplexed packet is equal or greater than 128 bytes, additional bytes are added. The first bit of each of the added bytes is the **LXT**. If **LXT** is set to 1, it means that there is an additional byte for expressing the length. This allows to multiplex packets of any length.

- **Protocol (8 bits)** is the *Protocol* field of the multiplexed packet, according to IANA "Assigned Internet Protocol Numbers". It is only included when the **SPB** of the First Multiplexing header is 0.

This is a Wireshark screenshot showing two multiplexed IP packets (*tun mode*), over IP protocol (*network mode*):

The screenshot shows a Wireshark capture of two multiplexed IP packets. The packet list at the top shows packets 4 through 15. Packet 9 is selected, showing details for Ethernet II, Internet Protocol Version 4, and Simplemux packet/frame multiplexer. The packet bytes pane at the bottom shows the raw data of the selected packet.

No.	Time	Time delta	Source	Destination	Protocol	Length	Info
4	1.917225	0.000417	192.168.129.134	34.107.221.82	TCP	54	58044 →
5	1.917437	0.000212	34.107.221.82	192.168.129.134	TCP	60	[TCP AC
6	1.917819	0.000382	34.107.221.82	192.168.129.134	TCP	60	[TCP AC
7	3.126242	1.208423	192.168.129.129	204.11.201.12	NTP	90	NTP Ver
8	3.332636	0.206394	204.11.201.12	192.168.129.129	NTP	90	NTP Ver
9	8.028835	4.696199	192.168.100.1	192.168.100.2	ICMP	206	Echo (p
10	8.029789	0.000954	192.168.100.2	192.168.100.1	ICMP	121	Echo (p
11	8.030004	0.000215	192.168.100.2	192.168.100.1	ICMP	121	Echo (p
12	8.199137	0.169133	VMware_27:a0:8f	VMware_ff:9f:ac	ARP	60	Who has
13	8.199162	0.000025	VMware_ff:9f:ac	VMware_27:a0:8f	ARP	60	192.168
14	10.048969	1.849807	192.168.100.1	192.168.100.2	ICMP	206	Echo (p
15	10.049817	0.000848	192.168.100.2	192.168.100.1	ICMP	121	Echo (p

Frame 9: 206 bytes on wire (1648 bits), 206 bytes captured (1648 bits) on Ethernet II, Src: VMware_a5:cb:6d (00:0c:29:a5:cb:6d), Dst: VMware_27:a0:8f (00:50:56:27:a0:8f)

Internet Protocol Version 4, Src: 192.168.129.134, Dst: 192.168.129.129

Simplemux packet/frame multiplexer

- 1... .. = Single Protocol Bit: All the multiplexed packets/frames belong to the same protocol (1)
- 1... .. = Length extension: Packet/frame length in 13 bits (1)
- ..00 0000 .101 0100 = Length: 84
- 0... .. = Second length extension: No more bytes for length (0)
- Protocol: IPv4 (4)
- Simplemux payload: 45000054188240004001d8d2c0a86401c0a8640208006e61e8400015ea28e6100000000...
- Internet Protocol Version 4, Src: 192.168.100.1, Dst: 192.168.100.2
- Internet Control Message Protocol

Simplemux packet/frame multiplexer

- 0... .. = Length extension: Packet/frame length in 7 bits (0)
- .101 0100 = Length: 84
- Simplemux payload: 4500005418c440004001d890c0a86401c0a864020800ae011e8400025fa28e6100000000...
- Internet Protocol Version 4, Src: 192.168.100.1, Dst: 192.168.100.2
- Internet Control Message Protocol

0000 00 50 56 27 a0 8f 00 0c 29 a5 cb 6d 08 00 45 00 PV.....)~m~E-
0010 00 c0 04 d2 00 00 40 fd f0 16 c0 a8 81 86 c0 a8@:.....
0020 81 81 c0 54 04 45 00 00 54 18 82 40 00 40 01 d8 ...T.E..T..@..
0030 d2 c0 a8 64 01 c0 a8 64 02 08 00 6e 66 1e 84 00 ...d..d...nf...
0040 01 5e a2 8e 61 00 00 00 00 b1 3d 0e 00 00 00 00 ^..a.....=
0050 00 10 11 12 13 14 15 16 17 18 19 1a 1b 1c 1d 1e .!~\$%&'()*+,-./01234567TE..T..
0060 1f 20 21 22 23 24 25 26 27 28 29 2a 2b 2c 2d 2e @:~.....d...d..
0070 2f 30 31 32 33 34 35 36 37 54 45 00 00 54 18 c4a...p..
0080 40 00 40 01 d8 90 c0 a8 64 01 c0 a8 64 02 08 00!~\$%&'()
0090 ae 01 1e 84 00 02 5f a2 8e 61 00 00 00 70 a1+,-./01 234567
00a0 0e 00 00 00 00 00 10 11 12 13 14 15 16 17 18 19
00b0 1a 1b 1c 1d 1e 1f 20 21 22 23 24 25 26 27 28 29
00c0 2a 2b 2c 2d 2e 2f 30 31 32 33 34 35 36 37

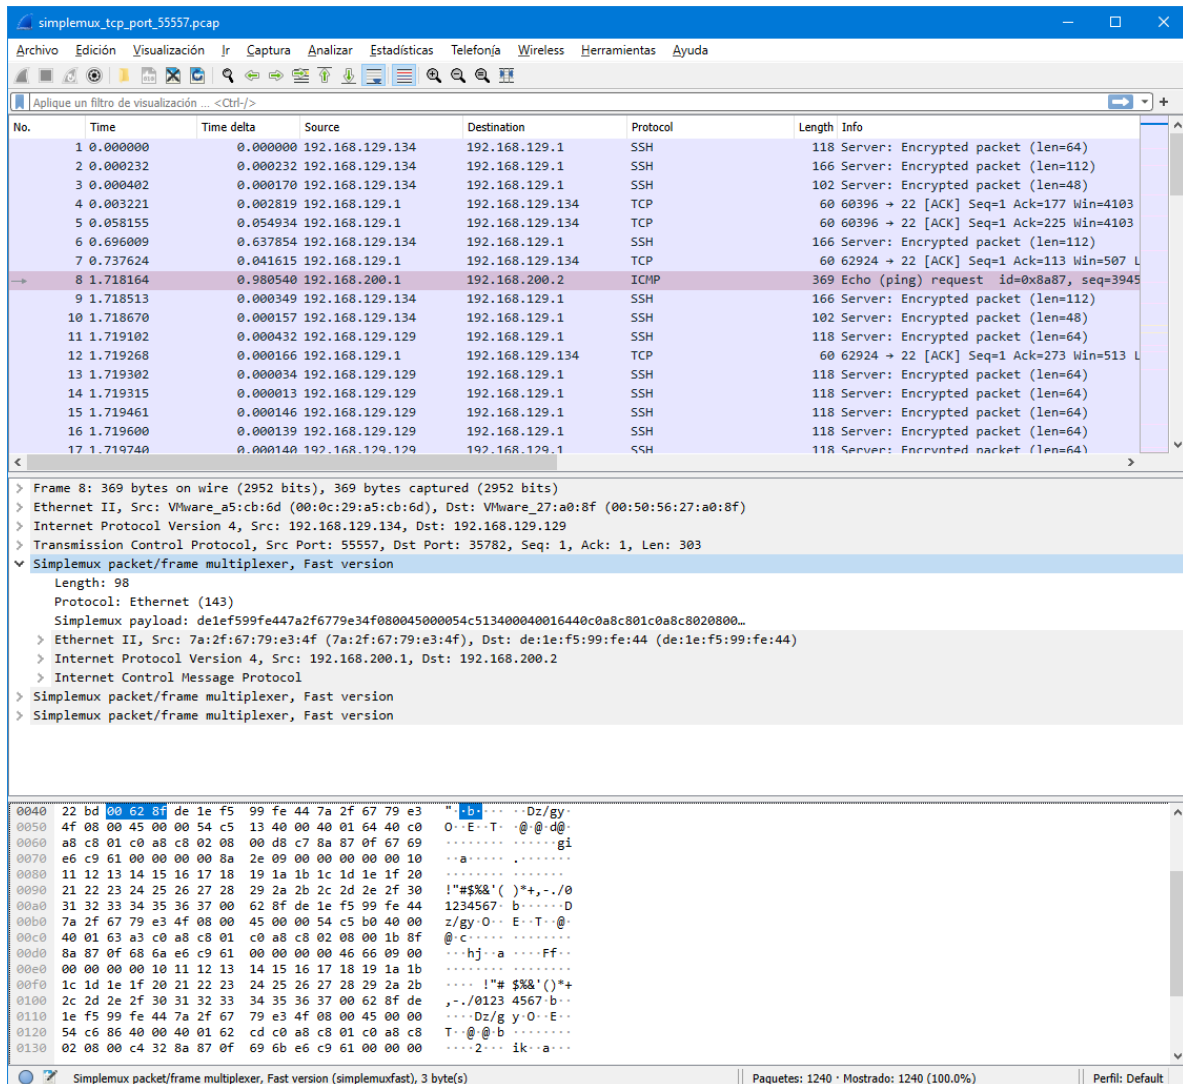
Internet Protocol Version 4 (ip), 20 byte(s) Paquetes: 56 · Mostrado: 56 (100.0%) Perfil: Default

Fast flavor

All the separators have the same structure:

- **Length (LEN, 16 bits).** This is the length of the multiplexed packet (in bytes), not including the *Length* field.
- **Protocol (8 bits)** is the *Protocol* field of the multiplexed packet, according to IANA "Assigned Internet Protocol Numbers".


This is a Wireshark screenshot showing three multiplexed ethernet frames (tap *tunnel mode*), over TCP protocol (tcp *mode*):



Multiplexing policies

Four different conditions can be used and combined for triggering the sending of a multiplexed packet (in the figures, the triggering moment is expressed by red lines):

- **number of packets:** a number of packets have arrived at the multiplexer.



- **size:** two different options apply:

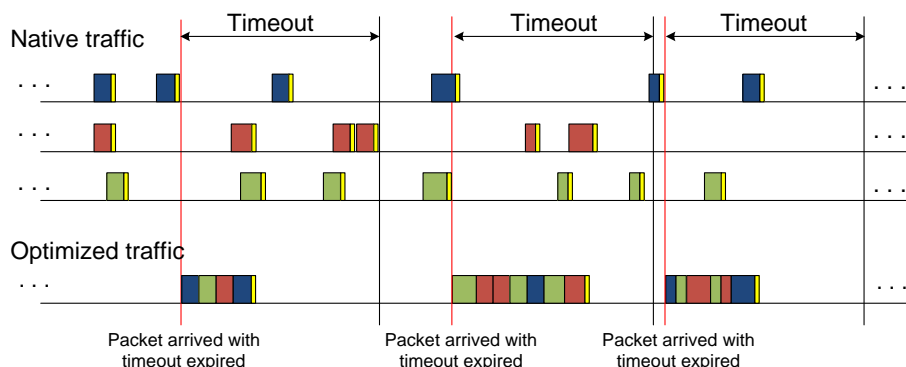
- the size of the multiplexed packet has exceeded the size threshold specified by the user, but not the MTU. In this case, a packet is sent, and a new period is started with the buffer empty.
- the size of the multiplexed packet has exceeded the MTU (and the size threshold consequently). In this case, a packet is sent without the last one. A new period is started, and the last arrived packet is stored for the next period.

If you want to specify an MTU different from the one of the local interface, you can use the `-m` option. You may use other tools for getting the MTU of a network path. For example with the command³:

```
# tracepath 192.168.137.3 | grep Resume | cut -c 19-22
```

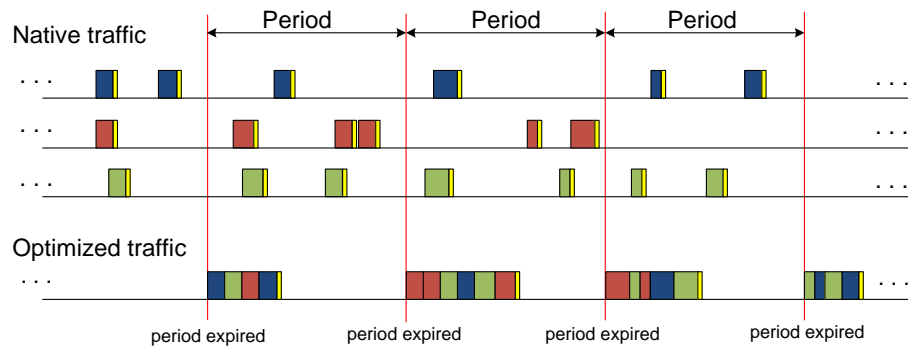
you will obtain the MTU of the path to 192.168.137.3.

- **timeout:** a packet arrives, and a timeout since the last sending has expired.



- **period:** an active waiting is performed, and a multiplexed packet including all the packets arrived during a period is sent.

³ <http://packetlife.net/blog/2008/aug/18/path-mtu-discovery/>



More than one condition can be set at the same time. Please note that if ($\text{timeout} > \text{period}$), then the timeout has no effect. Note that **only period policy guarantees an upper bound for the multiplexing delay.**

Simplemux is symmetric, i.e. both machines may act as ingress and egress simultaneously. However, different policies can be established at each of the optimizers, e.g. in one side you can send a multiplexed packet every two native ones, and in the other side you can set a timeout.

What will you find in this package

You will find two different source files:

<code>simplemux.c</code>	The complete version of Simplemux
<code>simplemux_multiplexing_delay.pl</code>	A perl script for reading the traces
<code>simplemux_throughput_pps.pl</code>	A perl script for reading the traces
<code>driveGnuPlotStreams.pl</code>	A perl script for creating real-time graphs
<code>simplemux.lua</code>	A dissector for Simplemux , <i>normal</i> flavor
<code>simplemuxfast.lua</code>	A dissector for Simplemux , <i>fast</i> flavor

Required tools

ROHC

First, you have to install rohc 1.7.0 from <https://rohc-lib.org/>. Download and uncompress the ROHC package.

```
cd rohc-1.7.0
./configure --prefix=/usr
make all
make check
make install
```

Compiling simplemux

```
gcc -o simplemux -g -Wall $(pkg-config rohc --cflags) simplemux.c $(pkg-config rohc --libs)
```

Usage of simplemux

```
./simplemux -i <ifacename> -e <ifacename> -c <peerIP> -M <'network' or 'udp' or 'tcpclient' or 'tcpserver'> [-T 'tun' or 'tap'] [-p <port>] [-d <debug_level>] [-r <ROHC_option>] [-n <num_mux_tun>] [-m <MTU>] [-b <num_bytes_threshold>] [-t <timeout (microsec)>] [-P <period (microsec)>] [-l <log file name>] [-L]
```

```
./simplemux -h
```

```
-i <ifacename>: Name of tun/tap interface to be used for capturing native packets (mandatory)
-e <ifacename>: Name of local interface which IP will be used for reception of muxed packets, i.e., the tunnel local end (mandatory)
-c <peerIP>: specify peer destination IP address, i.e. the tunnel remote end (mandatory)
-M <mode>: 'network' or 'udp' or 'tcpclient' or 'tcpserver' mode (mandatory)
-T <tunnel mode>: 'tun' (default) or 'tap' mode
-f: Fast mode (compression rate is lower, but it is faster). Compulsory for TCP mode
-p <port>: port to listen on, and to connect to (default 5555)
-d <debug_level>: Debug level. 0:no debug; 1:minimum debug; 2:medium debug; 3:maximum debug (incl. ROHC)
-r <ROHC_option>: 0:no ROHC; 1:Unidirectional; 2: Bidirectional Optimistic; 3: Bidirectional Reliable (not available yet)
-n <num_mux_tun>: number of packets received, to be sent to the network at the same time, default 1, max 100
-m <MTU>: Maximum Transmission Unit of the network path (by default the one of the local interface is taken)
-b <num_bytes_threshold>: size threshold (bytes) to trigger the departure of packets (default MTU-28 in transport mode and MTU-20 in network mode)
```

-t <timeout (microsec)>: timeout (in usec) to trigger the departure of packets
 -P <period (microsec)>: period (in usec) to trigger the departure of packets. If (timeout < period) then the timeout has no effect
 -l <log file name>: log file name. Use 'stdout' if you want the log data in standard output
 -L: use default log file name (day and hour Y-m-d_H.M.S)
 -h: prints this help text

If you have to use the same local interface more than once

It may happen that you have to create more than one tunnel using the same local interface. In that case, you may obtain a message **Is already in use**.

For example, if you run these two commands in the same machine, you may obtain this error message:

```
./simplemux -i tun0 -e wlan0 -M network -T tun -c 10.1.10.4
./simplemux -i tun1 -e wlan0 -M network -T tun -c 10.1.10.6
```

To avoid the problem, you can use an alias (see e.g. <https://www.cyberciti.biz/faq/linux-creating-or-adding-new-network-alias-to-a-network-card-nic/>). So you can do `ifconfig wlan0:0 x.y.z.t. up`, and then you can use `wlan0` in one case, and `wlan0:0` in the other.

Format of the Simplemux traces

Using the options `-l log file name` or `-L`, you can obtain a text file with traces. This is the format of these traces:

timestamp	event	type	size	sequence number	from/to	IP	port	number of packets	triggering event(s)		
%"PRIu64"	text	text	%i	%lu	text	%s	%d	%i	text		
microseconds	rec	native	packet size in bytes	sequence number	-	-	-	-	-		
		muxed			from	ingress IP address	port (only in transport mode)				
		ROHC_feedback									
	sent	muxed			to	egress IP address	port (only in transport mode)		number	numpacket_limit	
										size_limit	
		demuxed								timeout	
	error	native			-	-	-	-	-		
		bad_separator			-	-	-	-	-		
		demux_bad_length			-	-	-	-	-		
		decomp_failed			-	-	-	-	-		
		comp_failed			-	-	-	-	-		
drop	too_long	to	egress IP address	port	number	-					

	drop	no_ROHC_mode			-	-	-	-	-
--	------	--------------	--	--	---	---	---	---	---

- **timestamp**: it is in microseconds. It is obtained with the function `GetTimeStamp()`.
- **event** and **type**:
 - **rec**: a packet has been received:
 - **native**: a native packet has arrived to the ingress optimizer.
 - **muxed**: a multiplexed packet has arrived to the egress optimizer.
 - **ROHC_feedback**: a ROHC feedback-only packet has been received from the decompressor. It only contains ROHC feedback information, so there is nothing to decompress
 - **sent**: a packet has been sent
 - **muxed**: the ingress optimizer has sent a multiplexed packet.
 - **demuxed**: the egress optimizer has demuxed a native packet and sent it to its destination.
 - **forward**: when a packet arrives to the egress with a port different to the one where the optimization is being deployed, it is just forwarded to the network.
 - **error**:
 - **bad_separator**: the Simplemux header before the packet is not well constructed.
 - **demux_bad_length**: the length of the packet expressed in the Simplemux header is excessive (the multiplexed packet would finish after the end of the global packet).
 - **decomp_failed**: ROHC decompression failed.
 - **comp_failed**: ROHC compression failed.
 - **drop**:
 - **no_ROHC_mode**: a ROHC packet has been received, but the decompressor is not in ROHC mode.
- **size**: it expresses (in bytes) the size of the packet. If it is a **muxed** one, it is the global size of the packet (including IP header). If it is a **native** or **demuxed** one, it is the size of the original (native) packet.
- **sequence number**: it is a sequence number generated internally by the program. Two different sequences are generated: one for received packets and other one for sent packets.
- **IP**: it is the IP address of the peer Simplemux optimizer.
- **port**: it is the destination port of the packet.
- **number of packets**: it is the number of packets included in a multiplexed bundle.
- **triggering event(s)**: it is the cause (more than one may appear) of the triggering of the multiplexed bundle:
 - **numpacket_limit**: the limit of the number of packets has been reached.
 - **size_limit**: the maximum size has been reached.
 - **timeout**: a packet has arrived once the timeout had expired.
 - **period**: the period has expired.
 - **MTU**: the MTU has been reached.

Trace examples

In the ingress optimizer you may obtain:

```
1417693720928101  rec  native  63  1505
1417693720931540  rec  native  65  1506
1417693720931643  rec  native  52  1507
1417693720936101  rec  native  48  1508
1417693720936210  rec  native  53  1509
1417693720936286  rec  native  67  1510
1417693720937162  rec  native  57  1511
1417693720938081  sent  muxed   237 1511 to 192.168.137.4 55555 7 period
```

This means that 7 native packets (length 63, 65, ... 57, and sequence numbers 1505 to 1511) have been received, and finally the period has expired, so they have been sent together to the egress Simplemux optimizer at 192.168.137.4, port 55555.

In the egress optimizer you may obtain:

```
1417693720922848  rec  muxed   237 210  from 192.168.0.5 55555
1417693720922983  sent  demuxed  63 210
1417693720923108  sent  demuxed  65 210
1417693720923186  sent  demuxed  52 210
1417693720923254  sent  demuxed  48 210
1417693720923330  sent  demuxed  53 210
1417693720923425  sent  demuxed  67 210
1417693720923545  sent  demuxed  57 210
```

This means that a multiplexed packet (sequence number 210) has been received from the ingress optimizer 192.168.0.5 with port 55555, and it has been demuxed, resulting into 7 different packets of lengths 63, 65, ... 57.

Scripts for calculating compression statistics

This package includes the next Perl scripts:

Calculate throughput and packets per second

`simplemux_throughput_pps.pl`

It is able to calculate the throughput and the packet-per-second rate, from a Simplemux output trace. The result is in three columns:

tick_end_time(us)	throughput(bps)	packets_per_second
1000000	488144	763
2000000	490504	759
3000000	475576	749
4000000	483672	760
5000000	481784	758
6000000	487112	762
7000000	486824	760
8000000	488792	765
9000000	483528	761
10000000	486360	760

Usage:

```
$perl simplemux_throughput_pps.pl <trace file> <tick(us)> <event> <type> <peer IP> <port>
```

Examples:

```
# $ perl simplemux_throughput_pps.pl tracefile.txt 1000000 rec native all all
# $ perl simplemux_throughput_pps.pl log_simplemux 1000000 rec muxed all all
# $ perl simplemux_throughput_pps.pl log_simplemux 1000000 rec muxed 192.168.0.5 55555
# $ perl simplemux_throughput_pps.pl log_simplemux 1000000 sent demuxed
```

Calculate the multiplexing delay of each packet

Multiplexing delay is the time each packet is stopped in the multiplexer, i.e. the interval between its arrival as native packet and its departure inside a multiplexed packet.

`simplemux_multiplexing_delay.pl`

It is able to calculate the multiplexing delay of each packet, from a Simplemux output trace. The result is an output file in two columns:

```
packet_id    multiplexing_delay(us)
1           5279
2           1693
3           1202
4            507
5          10036
6           8471
7           6974
8           5588
9           1143
10          10435
11           8935
12           7522
13           5981
14          4520
15          3011
...
```

And other results are shown in stdout:

```
total native packets:      6661
Average multiplexing delay: 5222.47680528449 us
stdev of the multiplexing delay: 3425.575192789 us
```

Usage:

```
$ perl simplemux_multiplexing_delay.pl <trace file> <output file>
```

Scripts for drawing the instantaneous throughput and pps

You can make Simplemux generate real-time graphs of the throughput and the amount of packets per second. For that aim, you have to use pipes in order to combine two perl scripts.

You will need to install the `gnuplot-x11` Linux package.

- Send the log of simplemux to `stdout`, using `-l stdout` option.
- Use the script `simplemux_throughput_pps_live.pl` to generate a summary every e.g. 10 ms of packets coming from (or going to) 192.168.0.5 using port 55555.

The output is something like this:

```
0:492800
1:532000
2:700
3:700
0:492800
1:532000
2:700
3:700
```

Where each row represents a value of

0: native throughput

- 1: multiplexed throughput
- 2: native pps
- 3: multiplexed pps

- Use the script `driveGnuPlotStreams.pl`, by Andreas Bernauer⁴ to represent different graphs.

Some examples:

The next command presents two windows, each one with two graphs of 300 samples width, titled “native”, “muxed”, “ppsnat” and “ppsmux”

```
$ ./simplemux -i tun0 -e eth0 -c 192.168.0.5 -M udp -T tun -d 0 -r 2 -l stdout | perl
simplemux_throughput_pps_live.pl 10000 192.168.0.5 55555 |
perl ./driveGnuPlotStreams.pl 4 2 300 300 0 1000000 0 200 500x300+0+0 500x300+0+0 'native' 'muxed'
'ppsnat' 'ppsmux' 0 0 1 1
```

The next command presents one window with two graphs of 300 samples width, titled “native” and “muxed”

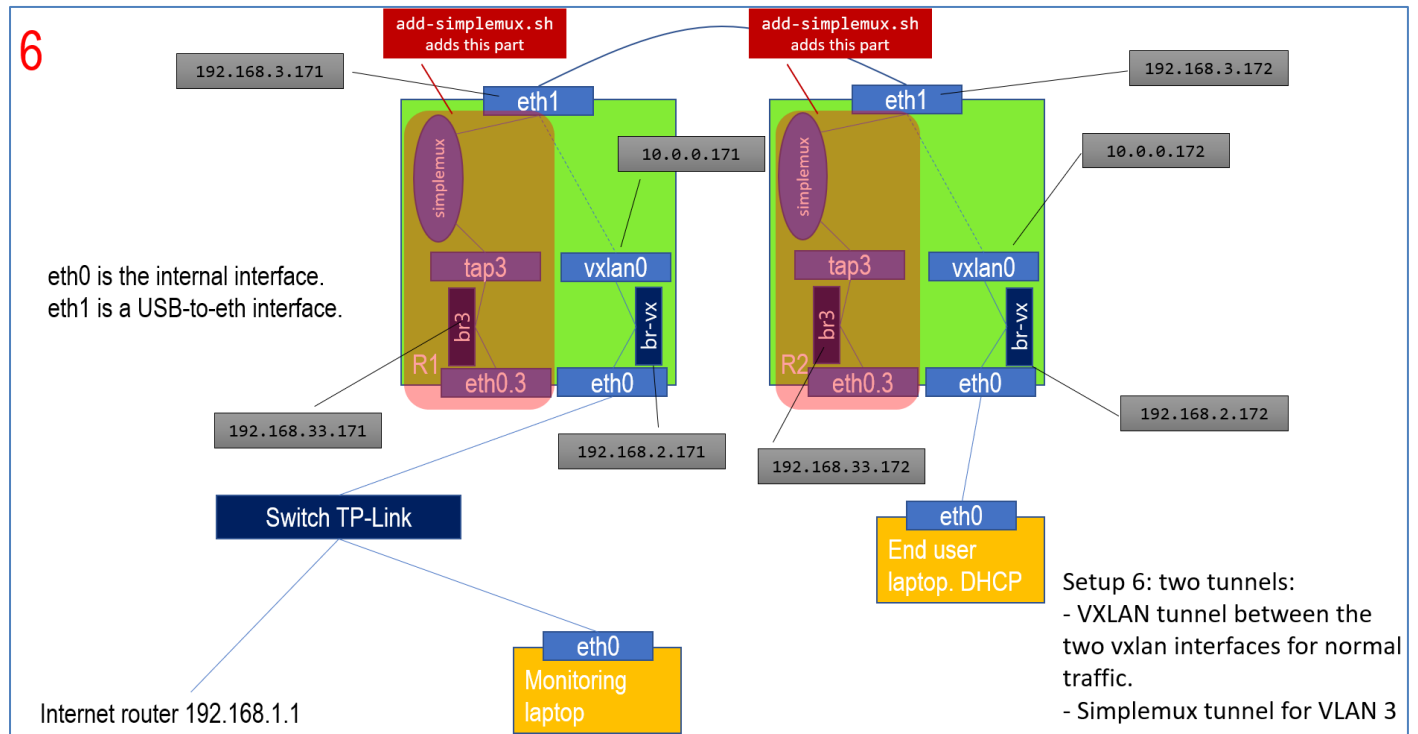
```
./simplemux -i tun0 -e eth0 -c 192.168.0.5 -M udp -T tun -d 0 -r 2 -l stdout | perl
simplemux_throughput_pps_live.pl 10000 192.168.0.5 | perl ./driveGnuPlotStreams.pl 2 1 300 0 1000000
500x300+0+0 'native' 'muxed' 0 0
```

⁴ <http://www.lysium.de/blog/index.php/?archives/234-Plotting-data-with-gnuplot-in-real-time.html>

Scenario 1: two Raspberries

Using tc qdisc and tc filter, you can build this scheme between two machines:

Note: the “right” part of each machine includes a VXLAN tunnel, but this is not important here.



With this script, you create the “red” part in the left machine:

```
# raspberry 171

# add the interface eth0.3, part of VLAN 3
ip link add link eth0 name eth0.3 type vlan id 3

# add the interface tap3
ip tuntap add dev tap3 mode tap user root

# set the interface up
ip link set tap3 up

# add a bridge where we will put eth0.3 and tap3
ip link add br3 type bridge

# add tap3 interface to the bridge
ip link set tap3 master br3

# add an IP address to tap3 (probably not needed)
ifconfig tap3 10.0.3.171 netmask 255.255.255.0

# add an IP address to br3 (useful for testing that the setup is correct)
ifconfig br3 192.168.33.171

# add eth0.3 to the bridge
ip link set eth0.3 master br3

# set the bridge up
ip link set br3 up

# create a tunnel to the other side

# using UDP:
#./simplemux/simplemux -i tap3 -e eth1 -M udp -T tap -c 192.168.3.172 -d 2

# using TCP server:
```

```
./simplemux/simplemux -i tap3 -e eth1 -M tcpserver -T tap -c 192.168.3.172 -d 3 -n 1 -f
```

With this script you can add the red part in the right machine:

```
# raspberry 172
# add the interface eth0.3, part of VLAN 3
ip link add link eth0 name eth0.3 type vlan id 3

# add the interface tap3
ip tuntap add dev tap3 mode tap user root

# set the interface up
ip link set tap3 up

# add a bridge where we will put eth0.3 and tap3
ip link add br3 type bridge

# add tap3 interface to the bridge
ip link set tap3 master br3

# add an IP address to tap3 (probably not needed)
ifconfig tap3 10.0.3.172 netmask 255.255.255.0

# add an IP address to br3 (useful for testing that the setup is correct)
ifconfig br3 192.168.33.172

# add eth0.3 to the bridge
ip link set eth0.3 master br3

# set the bridge up
ip link set br3 up

# create a tunnel to the other side
# using UDP:
./simplemux/simplemux -i tap3 -e eth1 -M udp -T tap -c 192.168.3.171 -d 2

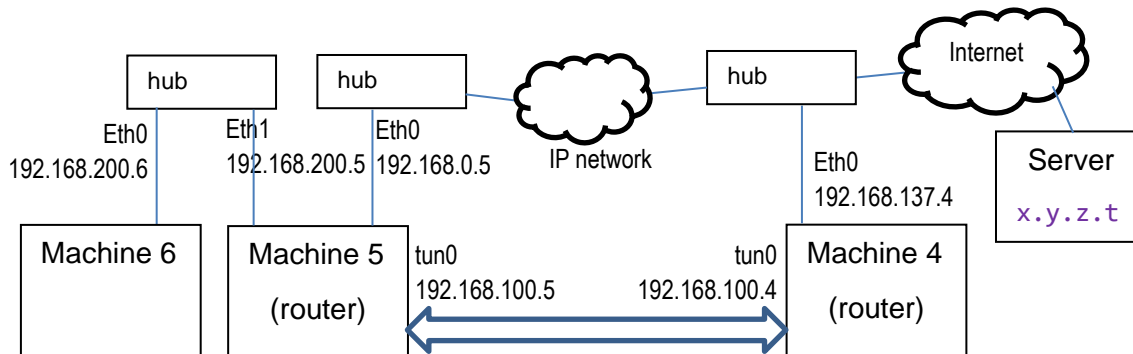
# using TCP, client:
./simplemux/simplemux -i tap3 -e eth1 -M tcpclient -T tap -c 192.168.3.171 -d 2 -n 1 -f
```

You can now ping from Raspberry 1 to 192.168.33.172, and you will see if the tunnel works.

Scenario 2: three machines

This is the setup:

Machine 6 is the source. Machine 5 and Machine 4 are the two optimizers. Server `x.y.z.t` is the destination.



Create a tun interface in machine 4

```
# ip tuntap add dev tun0 mode tun user root5
```

(`#openvpn --mktun --dev tun0 --user root` will work in OpenWrt and also in other Linux distributions⁶).

```
# ip link set tun0 up (or ifconfig tun0 up for e.g. OpenWRT)
```

If you want to add an IP address to the `tun0` interface, use:

```
# ip addr add 192.168.100.4/24 dev tun0
```

If you do not need an IP address, you can omit the previous command.

Create a tun interface in machine 5

```
# ip tuntap add dev tun0 mode tun user root
```

(`#openvpn --mktun --dev tun0 --user root` will also work)

```
# ip link set tun0 up
```

```
# ip addr add 192.168.100.5/24 dev tun0
```

Establish the simplemux tunnel between machine 4 and machine 5

In machine4:

```
# ./simplemux -i tun0 -e eth0 -M udp -T tun -c 192.168.0.5
```

In machine5:

```
# ./simplemux -i tun0 -e eth0 -M udp -T tun -c 192.168.137.4
```

Now you can ping from machine 5 or machine 6, to machine 4:

```
$ ping 192.168.100.4
```

The ping arrives to the `tun0` interface of machine 5, goes to machine 4 through the tunnel and is returned to machine 6 through the tunnel.

⁵ For removing the interface use `ip tuntap del dev tun0 mode tun`

⁶ Openvpn is used to create and destroy tun/tap devices. In Debian you can install it this way:

```
# apt-get install openvpn
```

In OpenWRT you will not be able to run `ip tuntap`, so you should install `openvpn` with:

```
# opkg install openvpn-noss1 (do opkg update before).
```

How to steer traffic from Machine 6 to server x.y.z.t through the tunnel

Steer the desired traffic to the tunnel using *ip rule* and *iptables*

The idea of **simplemux** is that it does not run at endpoints, but on some “optimizing” machines in the network. Therefore, you have to define policies to steer the flows of interest, in order to make them go through the TUN interface of the ingress (machine 5). This can be done with *ip rule* and *iptables*.

Following with the example:

In Machine 5, add a rule that makes the kernel route packets marked with 2 through *table 3*:

```
# ip rule add fwmark 2 table 3
```

In Machine 5, add a new route for *table 3*:

```
# ip route add default dev tun0 table 3 7
# ip route flush cache
```

If you show the routes of *table 3*

```
# ip route show table 3
```

Then you should obtain this:

```
default via 192.168.100.5 dev tun0
```

And now you can use *iptables* in order to mark certain packets as 2 if they have a certain destination IP, or a port number.

Adding *iptables* entries:

All packets with destination IP address x.y.z.t

```
iptables -t mangle -A PREROUTING -p udp -d x.y.z.t -j MARK --set-mark 2
```

All packets with destination UDP port 8999

```
iptables -t mangle -A PREROUTING -p udp --dport 8999 -j MARK --set-mark 2
```

All packets with destination TCP port 44172

```
iptables -t mangle -A PREROUTING -p tcp --dport 44172 -j MARK --set-mark 2
```

Remove the table rule

```
iptables -t mangle -D PREROUTING -p tcp --dport 44172 -j MARK --set-mark 2
```

Show the table

```
iptables -t mangle -L
```

Examples implementing different policies

Set a period of 50 ms

```
./simplemux -i tun0 -e eth0 -M udp -T tun -c 192.168.0.5 -P 50000
```

Send a multiplexed packet every 2 packets, use ROHC Bidirectional Optimistic

```
./simplemux -i tun0 -e eth0 -M udp -T tun -c 192.168.0.5 -n 2 -r 2
```

Send a multiplexed packet if the size of the multiplexed bundle is 400 bytes

```
./simplemux -i tun0 -e eth0 -M udp -T tun -c 192.168.0.5 -b 400
```

⁷ If you have set an IP address in the *tun0* interface, this command should also work:

```
# ip route add default via 192.168.100.5 table 3
```

Send a timeout of 50ms, and a period of 100 ms (to set an upper bound on the added delay), use ROHC Unidirectional

```
./simplemux -i tun0 -e eth0 -M udp -T tun -c 192.168.0.5 -t 50000 -P 100000 -r 1
```

Credits

The author of **Simplemux** is Jose Saldana. **Simplemux** has been written for research purposes, so if you find it useful, I would appreciate that you send a message sharing your experiences, and your improvement suggestions.

The software is released under the **GNU General Public License**, Version 3, 29 June 2007.

Thanks to Didier Barvaux for his ROHC implementation.

Thanks to Davide Brini for his simuletun program.

If you have some improvement suggestions, do not hesitate to contact me.

The extension from Simplemux to Simplemux (October 2021) has been done by Jose Saldana, working at [CIRCE Foundation](#), as a part of the [H2020 FARCROSS project](#). This project has received funding from the European Union's Horizon 2020 research and innovation programme under grant agreement No 864274.