Simplemux Readme file

About Simplemux

There are some situations in which multiplexing a number of small packets into a bigger one is desirable. For example, a number of small packets can be sent together between a pair of machines if they share a common network path. Thus, the traffic profile can be shifted from small to larger packets, reducing the network overhead and the number of packets per second to be managed by intermediate routers.

Simplemux is a generic multiplexing protocol, described in draft-saldana-tsvwg-simplemux. It is able to encapsulate a number of packets belonging to different protocols into a single packet. It includes the "Protocol" field on each multiplexing   header, thus allowing the inclusion of a number of packets belonging to different protocols in a packet of another protocol (Fig. 1). The size of the multiplexing headers is kept very low (it may be a single byte when multiplexing small packets) in order to reduce the overhead.


Fig. 1. Example of a Simplemux packet including packets of different protocols

Simplemux is designed to optimize together a number of flows sharing a common network path or segment (Fig. 2). Optimization in the end host is (in principle) not useful, since a number of small-packet flows departing from the same host are unusual. The multiplexing is performed between a pair of machines called ingress-optimizer and egress-optimizer.


Fig. 2. Scheme of the optimization

The present code is an implementation of **Simplemux**, using IP/UDP as the *multiplexing protocol*, and IP or ROHC (RFC 3095) as the *multiplexed protocol*, as illustrated in Fig. 3. Thus, it is able to perform traffic optimization, combining multiplexing with header compression. This may result on significant bandwidth savings and pps reductions for small-packet flows (e.g. VoIP, online games).

Simplemux can run in two modes:

- **Network mode**: the multiplexed packet is sent in an IP datagram using Protocol Number 253 (according to IANA, this number can be used for experimentation and testing[1]).

- **Transport mode**: the multiplexed packet is sent in an IP/UDP datagram. In this case, the protocol number in the IP header is that of UDP (17) and both ends have to agree on a UDP port.

**Native traffic: Five** IPv4/UDP/RTP VoIP packets with two samples of 10 bytes

Optimized traffic (network mode): One IPv4 simplemux Packet including **five** RTP packets

saving

Optimized traffic (transport mode): One IPv4 simplemux Packet including **five** RTP packets

saving

Native traffic headers:

	IPv4 header: 20 bytes
	UDP header: 8 bytes
	RTP header: 12 bytes
	Payload

Optimized traffic headers:

| | Tunnel IP header: 20 bytes |
	Tunnel UDP header: 8 bytes
	Simplemux separator: 1-3 bytes
	ROHC Compressed header (IP/UDP/RTP)

Fig. 3. Optimization of five RTP packets of VoIP

Simplemux can be used as an option at the multiplexing layer of TCM[2], a protocol combining **T**unneling, **C**ompressing and **M**ultiplexing for the optimization of small-packet flows. TCM may use of a number of different standard algorithms for header compression, multiplexing and tunneling, combined in a similar way to RFC 4170, as proposed in draft-saldana-tsvwg-tcmtf.

The implementation is written in C for Linux. It compresses headers using an implementation of ROHC by Didier Barvaux (https://rohc-lib.org/).

Simplemux uses Linux TUN virtual interface.

---

[1] Protocol numbers, http://www.iana.org/assignments/protocol-numbers/protocol-numbers.xhtml

[2] Jose Saldana *et al, "Emerging Real-time Services: Optimizing Traffic by Smart Cooperation in the Network,"* IEEE Communications Magazine, Vol. 51, n. 11, pp 127-136, Nov. 2013. doi 10.1109/MCOM.2013.6658664

Header compression

Simplemux uses an implementation of ROHC by Didier Barvaux (https://rohc-lib.org/). It implements these ROHC modes:

- ROHC unidirectional.

- ROHC bidirectional optimistic

- ROHC bidirectional reliable is not yet implemented.

ROHC cannot be enabled in one of the peers and disabled in the other peer.

ROHC is able to compress this traffic flows:

- IP/UDP/RTP: If the UDP packets have the destination ports 1234, 36780, 33238, 5020, 5002, the compressor assumes that they are RTP.
- IP/UDP
- IP/TCP
- IP/ESP
- IP/UDP-Lite

Tunneling

In this **Simplemux** implementation, tunneling is performed in a very simple way:

- In **network mode**, the external IP header is the tunneling header. It uses Protocol number 253.

- In **transport mode**, an IP/UDP header is added before the first **Mux separator.** By default, the destination UDP port is 55555.

ROHC feedback information (when using ROHC Bidirectional modes) is sent in UDP packets using port 55556 by default.

Multiplexing

The **Mux separator** (see Fig. 4) has two different formats: one for the *First header* (the separator before the first packet included in the multiplexed bundle), and another one for *Non-first headers* (the rest of the separators).

First Simplemux header   Non-first Simplemux headers

Fig. 4. First and Non-first Simplemux headers (also known as separators)

Format of the First Simplemux header/separator

The format of the First Simplemux separator is illustrated in Fig. 5:

packet length < 64 bytes

| SPB | LXT | LEN (6 bits) | Protocol (8 bits) |

packet length ≥ 64 bytes

| SPB | LXT | LEN (14 bits) | Protocol (8 bits) |

Fig. 5. Fields of the First Simplemux header/separator

- **Single Protocol Bit (SPB, one bit)** only appears in the first Simplemux header. It is set to 1 if all the multiplexed packets belong to the same protocol (in this case, the "protocol" field will only appear in the first Simplemux header). It is set to 0 when each packet MAY belong to a different protocol.

- **Length Extension (LXT, one bit)** is 0 if the length of the first packet can be expressed in 6 bits, and 1 in other case.

- **Length (LEN, 6 or 14 bits)**: This is the length of the multiplexed packet in bytes not including the length field. If the length of the multiplexed packet is less than 64 bytes (less than or equal to 63 bytes), LXT is set to 0 and the 6 bits of the length field are the length of the multiplexed packet. If the length of the multiplexed packet is greater than 63 bytes, LXT is set to 1 and the 14 bits of the length field are the length of the multiplexed packet. The maximum length of a multiplexed packet is 16,383 bytes. Packets larger than 16,383 bytes will need to be sent in their native form. A Simplemux ingress is not required to multiplex all packets smaller than 16,383 bytes. It may choose to only multiplex packets smaller than a configurable size into a Simplemux multiplexed packet.

- **Protocol (8 bits)** is the Protocol field of the multiplexed packet, according to IANA "Assigned Internet Protocol Numbers".

Format of the Non-first Simplemux header/separator

If the SPB (Single Protocol Bit) of the First Simplemux header is set to 1, it means that all the multiplexed packets belong to the same protocol. In this case, the format is the one presented in Fig. 6:


Fig. 6. Fields of the Non-first Simplemux header/separator, when the SPB bit of the First header is 1

If the SPB of the First Simplemux header is set to 0, then each packet may belong to a different protocol, so the "Protocol" field is also included, as shown in Fig. 7:


Fig. 7. Fields of the Non-first Simplemux header/separator, when the SPB bit of the First header is 0

These are the fields:

- **Length Extension (LXT, one bit)** is 0 if the length of the first packet can be expressed in 7 bits, and 1 in other case.

- **Length (LEN, 7 or 15 bits)**: This is the length of the multiplexed packet in bytes not including the length field. If the length of the multiplexed packet is less than 128 bytes (less than or equal to 127 bytes), LXT is set to 0 and the 7 bits of the length field represent the length of the multiplexed packet. If the length of the multiplexed packet is greater than 127 bytes, LXT is set to 1 and the 15 bits of the length field are the length of the multiplexed packet. The maximum length of a multiplexed packet is 32,768 bytes.

Packets larger than 32,768 bytes will need to be sent in their native form. However, this will have to be reduced to 16,383 bytes taking into account that the maximum size of the First header is 14 bits. A Simplemux ingress is not required to multiplex all packets smaller than 32,768 bytes. It may choose to only multiplex packets smaller than a configurable size into a Simplemux multiplexed packet.

- **Protocol (8 bits)** is the Protocol field of the multiplexed packet, according to IANA "Assigned Internet Protocol Numbers". It is only included when the SPB of the First Multiplexing header is 0.

Multiplexing policies

Four different conditions can be used and combined for triggering the sending of a multiplexed packet (in the figures, the triggering moment is expressed by red lines):

- **number of packets**: a number of packets have arrived to the multiplexer.

- **size**: two different options apply:

  - the size of the multiplexed packet has exceeded the size threshold specified by the user, but not the MTU. In this case, a packet is sent and a new period is started with the buffer empty.

  - the size of the multiplexed packet has exceeded the MTU (and the size threshold consequently). In this case, a packet is sent without the last one. A new period is started, and the last arrived packet is stored for the next period.

If you want to specify an MTU different from the one of the local interface, you can use the '-m' option. You may use other tools for getting the MTU of a network path. For example with the command[3]:

```
# tracepath 192.168.137.3 | grep Resume | cut -c 19-22
```

you will obtain the MTU of the path to 192.168.137.3

---

[3] http://packetlife.net/blog/2008/aug/18/path-mtu-discovery/

- **timeout**: a packet arrives, and a timeout since the last sending has expired.

- **period**: an active waiting is performed, and a multiplexed packet including all the packets arrived during a period is sent.

More than one condition can be set at the same time. Please note that if (timeout < period), then the timeout has no effect. Note that **only period policy guarantees an upper bound for the multiplexing delay.**

Simplemux is symmetric, i.e. both machines may act as ingress and egress simultaneously. However, different policies can be established at each of the optimizers, e.g. in one side you can send a multiplexed packet every two native ones, and in the other side you can set a timeout.

What will you find in this package

You will find two different source files:

simplemux.c	The complete version of **Simplemux**
simplemux_multiplexing_delay.pl	A perl script for reading the traces
simplemux_throughput_pps.pl	A perl script for reading the traces

Required tools

ROHC (not required for simplemux-no-compress.c)

First, you have to install rohc 1.7.0 from https://rohc-lib.org/. Download and uncompress the ROHC package.

```
cd rohc-1.7.0
./configure --prefix=/usr
make all
make check
make install
```

Compiling simplemux

```
gcc -o simplemux -g -Wall $(pkg-config rohc --cflags) simplemux_vX.Y.c
$(pkg-config rohc --libs )
```

Usage of simplemux

```
./simplemux -i <ifacename> -e <ifacename> -c <peerIP> -M <N or T> [-p
<port>] [-d <debug_level>] [-r <ROHC_option>] [-n <num_mux_tun>] [-m
<MTU>] [-b <num_bytes_threshold>] [-t <timeout (microsec)>] [-P
<period (microsec)>] [-l <log file name>] [-L]


./simplemux -h


-i <ifacename>: Name of tun interface to use for capturing native
packets (mandatory)
-e <ifacename>: Name of local interface which IP will be used for
reception of muxed packets, i.e., the tunnel local end (mandatory)
-c <peerIP>: specify peer destination IP address, i.e. the tunnel
remote end (mandatory)
-M <mode>: Network(N) or Transport (T) mode (mandatory)
-p <port>: port to listen on, and to connect to (default 55555)
-d: outputs debug information while running. 0:no debug; 1:minimum
debug; 2:medium debug; 3:maximum debug (incl. ROHC)
-r: 0:no ROHC; 1:Unidirectional; 2: Bidirectional Optimistic; 3:
Bidirectional Reliable (not available yet)
-n: number of packets received, to be sent to the network at the same
time, default 1, max 100
-m: Maximum Transmission Unit of the network path (by default the one
of the local interface is taken)
-b: size threshold (bytes) to trigger the departure of packets
(default MTU-28 in transport mode and MTU-20 in network mode)
-t: timeout (in usec) to trigger the departure of packets
-P: period (in usec) to trigger the departure of packets. If ( timeout
< period ) then the timeout has no effect
-l: log file name
-L: use default log file name (day and hour Y-m-d_H.M.S)
-h: prints this help text
```

Format of the Simplemux traces

Using the options `-l log file name` or `-L` you can obtain a text file with traces. This is the format of these traces:

timestamp	event	type	size	sequence number	from/to	IP	port	number of packets	triggering event(s)
%"PRIu64"	text	text	%i	%lu	text	%s	%d	%i	text
microseconds	rec	native	packet size in bytes	sequence number	-	-	-	-	-
		muxed			from	ingress IP address	port	-	-
		ROHC_feedback							
	sent	muxed			to	egress IP address	port	number	numpacket_limit size_limit timeout period MTU
		demuxed			-	-	-	-	-
	forward	native			from	ingress IP address	port	-	-
	error	bad_separator			-	-	-	-	-
		demux_bad_length			-	-	-	-	-
		decomp_failed			-	-	-	-	-
		comp_failed			-	-	-	-	-
	drop	no_ROHC_mode			-	-	-	-	-

- `timestamp`: it is in microseconds. It is obtained with the function `GetTimeStamp()`.

- `event` and `type`:

  - `rec`: a packet has been received:

    - `native`: a native packet has arrived to the ingress optimizer.

    - `muxed`: a multiplexed packet has arrived to the egress optimizer.

    - `ROHC_feedback`: a ROHC feedback-only packet has been received from the decompressor. It only contains ROHC feedback information, so there is nothing to decompress

  - `sent`: a packet has been sent

    - `muxed`: the ingress optimizer has sent a multiplexed packet.

- `demuxed`: the egress optimizer has demuxed a native packet and sent it to its destination.

- `forward`: when a packet arrives to the egress with a port different to the one where the optimization is being deployed, it is just forwarded to the network.

- `error`:

  - `bad_separator`: the Simplemux header before the packet is not well constructed.

  - `demux_bad_length`: the length of the packet expressed in the Simplemux header is excessive (the multiplexed packet would finish after the end of the global packet).

  - `decomp_failed`: ROHC decompression failed.

  - `comp_failed`: ROHC compression failed.

- `drop`:

  - `no_ROHC_mode`: a ROHC packet has been received, but the decompressor is not in ROHC mode.

- `size`: it expresses (in bytes) the size of the packet. If it is a muxed one, it is the global size of the packet (including IP header). If it is a `native` or `demuxed` one, it is the size of the original (native) packet.

- `sequence number`: it is a sequence number generated internally by the program. Two different sequences are generated: one for received packets and other one for sent packets.

- `IP`: it is the IP address of the peer Simplemux optimizer.

- `port`: it is the destination port of the packet.

- `number of packets`: it is the number of packets included in a multiplexed bundle.

- `triggering event(s)`: it is the cause (more than one may appear) of the triggering of the multiplexed bundle:

  - `numpacket_limit`: the limit of the number of packets has been reached.

  - `size_limit`: the maximum size has been reached.

  - `timeout`: a packet has arrived once the timeout had expired.

  - `period`: the period has expired.

  - `MTU`: the MTU has been reached.

9

Trace examples

In the ingress optimizer you may obtain:

```
1417693720928101  rec   native   63  1505
1417693720931540  rec   native   65  1506
1417693720931643  rec   native   52  1507
1417693720936101  rec   native   48  1508
1417693720936210  rec   native   53  1509
1417693720936286  rec   native   67  1510
1417693720937162  rec   native   57  1511
1417693720938081  sent  muxed    237 1511  to    192.168.137.4  55555 7  period
```

This means that 7 native packets (length 63, 65, … 57, and sequence numbers 1505 to 1511) have been received, and finally the period has expired, so they have been sent together to the egress Simplemux optimizer at 192.168.137.4, port 55555.

In the egress optimizer you may obtain:

```
1417693720922848  rec   muxed    237 210   from  192.168.0.5   55555
1417693720922983  sent  demuxed  63  210
1417693720923108  sent  demuxed  65  210
1417693720923186  sent  demuxed  52  210
1417693720923254  sent  demuxed  48  210
1417693720923330  sent  demuxed  53  210
1417693720923425  sent  demuxed  67  210
1417693720923545  sent  demuxed  57  210
```

This means that a multiplexed packet (sequence number 210) has been received from the ingress optimizer 192.168.0.5 with port 55555, and it has been demuxed, resulting into 7 different packets of lengths 63, 65, … 57.

Scripts for calculating some statistics

This package includes the next Perl scripts:

Calculate throughput and packets per second

`simplemux_throughput_pps.pl`

It is able to calculate the throughput and the packet-per-second rate, from a Simplemux output trace. The result is in three columns:

```
tick_end_time(us)    throughput(bps)    packets_per_second
1000000              488144             763
2000000              490504             759
3000000              475576             749
4000000              483672             760
5000000              481784             758
6000000              487112             762
7000000              486824             760
8000000              488792             765
9000000              483528             761
10000000             486360             760
```

Usage:

```
$perl simplemux_throughput_pps.pl <trace file> <tick(us)> <event> <type> <peer IP> <port>
```

Examples:

```
# $ perl simplemux_throughput_pps.pl tracefile.txt 1000000 rec native all all
# $ perl simplemux_throughput_pps.pl log_simplemux 1000000 rec muxed all all
# $ perl simplemux_throughput_pps.pl log_simplemux 1000000 rec muxed 192.168.0.5 55555
# $ perl simplemux_throughput_pps.pl log_simplemux 1000000 sent demuxed
```

Calculate the multiplexing delay of each packet

Multiplexing delay is the time each packet is stopped in the multiplexer, i.e. the interval between its arrival as native packet and its departure inside a multiplexed packet.

```
simplemux_multiplexing_delay.pl
```

It is able to calculate the multiplexing delay of each packet, from a Simplemux output trace. The result is an output file in two columns:

```
packet_id multiplexing_delay(us)
1        5279
2        1693
3        1202
4        507
5        10036
6        8471
7        6974
8        5588
9        1143
10       10435
11       8935
12       7522
13       5981
14       4520
15       3011
...
```

And other results are shown in stdout:

```
total native packets: 6661
Average multiplexing delay:   5222.47680528449 us
stdev of the multiplexing delay:    3425.575192789 us
```
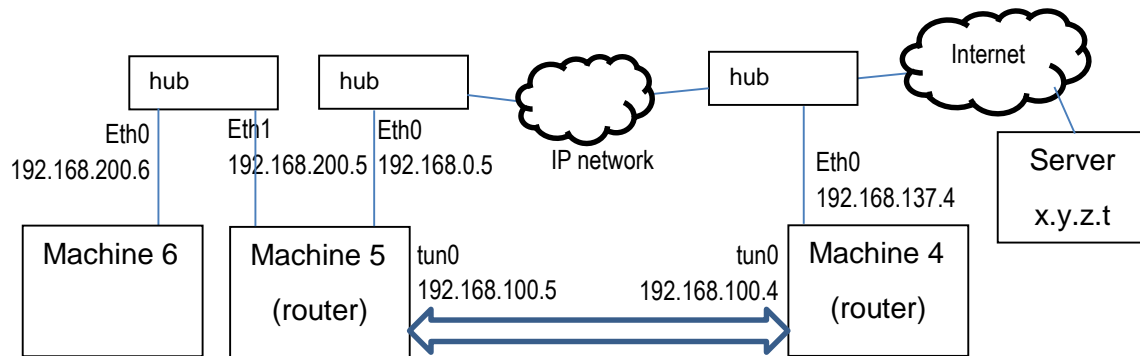
Usage:
```
$ perl simplemux_multiplexing_delay.pl <trace file> <output file>
```

Usage example with three machines

This is the setup:

Machine 6 is the source. Machine 5 and Machine 4 are the two optimizers. Server x.y.z.t is the destination.

Create a tun interface in machine 4

```
ip tuntap add dev tun0 mode tun user root⁴
```
(`openvpn --mktun --dev tun0 --user root` will also work[5]).

```
ip link set tun0 up
```
(or `ifconfig tun0 up` for e.g. OpenWRT)

If you want to add an IP address to the tun0 interface, use:

```
ip addr add 192.168.100.4/24 dev tun0
```

If you do not need an IP address, you can omit the previous command.

Create a tun interface in machine 5

```
ip tuntap add dev tun0 mode tun user root
```
(`openvpn --mktun --dev tun0 --user root` will also work)

```
ip link set tun0 up
```
```
ip addr add 192.168.100.5/24 dev tun0
```

Establish the simplemux tunnel between machine 4 and machine 5

In machine4:
```
# ./simplemux -i tun0 -e eth0 -M N -c 192.168.0.5
```

In machine5:
```
# ./simplemux -i tun0 -e eth0 -M N -c 192.168.137.4
```

Now you can ping from machine 5 or machine6 to machine 4:

```
$ ping 192.168.100.4
```

---

[4] For removing the interface use `ip tuntap del dev tun0 mode tun`

[5] Openvpn is used to create and destroy tun/tap devices. In Debian you can install it this way: `#apt-get install openvpn`

In OpenWRT you will not be able to run `ip tuntap`, so you should install openvpn with: `#opkg install openvpn-nossl`. (do opkg update before)

The ping arrives to the `tun0` interface of machine 5, goes to machine 4 through the tunnel and is returned to machine 6 through the tunnel.

How to steer traffic from Machine 6 to server x.y.z.t through the tunnel

The idea of **simplemux** is that it does not run on endpoints, but on some "optimizing" machines in the network. So you have to define policies to steer the flows of interest, in order to make them go through the TUN interface of the ingress (machine 5). This can be done with `iprules` and `iptables`.

Following with the example:

In Machine 5, add a rule that makes the kernel route packets marked with "2" through `table 3`:

```
# ip rule add fwmark 2 table 3
```

In Machine 5, add a new route for `table 3`:

```
# ip route add default dev tun0 table 3 ⁶
# ip route flush cache
```

If you show the routes of table 3

```
# ip route show table 3
```

Then you should obtain this:

```
        default via 192.168.100.5 dev tun0
```

And now you can use `iptables` in order to mark certain packets as "2" if they have a certain destination IP, or a port number.

Examples:

All packets with destination IP address x.y.z.t
```
iptables -t mangle -A PREROUTING -p udp -d x.y.z.t -j MARK --set-mark 2
```

All packets with destination UDP port 8999
```
iptables -t mangle -A PREROUTING -p udp --dport 8999 -j MARK --set-mark 2
```

All packets with destination TCP port 44172
```
iptables -t mangle -A PREROUTING -p tcp --dport 44172 -j MARK --set-mark 2
```

Remove the table rule
```
iptables -t mangle -D PREROUTING -p tcp --dport 44172 -j MARK --set-mark 2
```

Show the table
```
iptables -t mangle -L
```

---

⁶ If you have set an IP address in the tun0 interface, this command should also work:
```
# ip route add default via 192.168.100.5 table 3
```

Other examples implementing different policies

Set a period of 50 ms

```
./simplemux -i tun0 -e eth0 —M N -c 192.168.0.5 –P 50000
```

Send a multiplexed packet every 2 packets, use ROHC Bidirectional Optimistic

```
./simplemux -i tun0 -e eth0 —M N -c 192.168.0.5–n 2 –r 2
```

Send a multiplexed packet if the size of the multiplexed bundle is 400 bytes

```
./simplemux -i tun0 -e eth0 —M N -c 192.168.0.5 –b 400
```

Send a timeout of 50ms, and a period of 100 ms (to set an upper bound on the added delay), use ROHC Unidirectional

```
./simplemux -i tun0 -e eth0 —M N -c 192.168.0.5 –t 50000 –P 100000 –r 1
```

Credits

The author of **simplemux** is Jose Saldana (jsaldana at unizar.es). **Simplemux** has been written for research purposes, so if you find it useful, I would appreciate that you send a message sharing your experiences, and your improvement suggestions.

The software is released under the **GNU General Public License**, Version 3, 29 June 2007.

Thanks to Didier Barvaux for his ROHC implementation.

Thanks to Davide Brini for his simpletun program.

If you have some improvement suggestions, do not hesitate to contact me.

http://diec.unizar.es/~jsaldana/

Credits